

Moving Forward to C++11

Howard Hinnant

May 15, 2012

Outline

- The rvalue reference
- Move Semantics
- Factory Functions
- More rvalue ref rules
- “Perfect” forwarding


Outline

- The rvalue reference
- Move Semantics
- Factory Functions
- More rvalue ref rules
- “Perfect” forwarding

Motivation for Move

- In 2006, I wrote a benchmark to show off move semantics.
- It manipulated the unlikely data structure `vector<set<int>>`:
 - Return it from factory functions.
 - Manipulate it with algorithms.


Motivation for Move




construction
destroy

stable_sort


Motivation for Move




Motivation for Move




Motivation for Move




Motivation for Move




Motivation for Move



Motivation for Move



Motivation for Move



Motivation for Move

Motivation for Move

- With move semantics, `vector<set<int>>` does not have to be an unlikely data structure.

Motivation for Move

- With move semantics, `vector<set<int>>` does not have to be an unlikely data structure.
- Containers and algorithms can move around `set<int>` almost as cheaply as moving around an `int`.

Motivation for Move

- With move semantics, `vector<set<int>>` does not have to be an unlikely data structure.
- Containers and algorithms can move around `set<int>` almost as cheaply as moving around an `int`.
- And you can install move semantics in your “heavy” data structures.

Rvalue reference syntax

A&

Rvalue reference syntax

A&

- In C++03 we have the reference.

Rvalue reference syntax

A&

Rvalue reference syntax

A&

- In C++11 we renamed “reference” to “lvalue reference.”

Rvalue reference syntax

A&


A&&

- In C++11 we renamed “reference” to “lvalue reference.”
- And we introduce a new kind of reference called “rvalue reference.”

Expressions



Expressions




- In C++98/03 every expression is lvalue or rvalue.
- Expressions never have reference type.

Expressions




Expressions




- In C++11 we renamed rvalue to prvalue.

Expressions




- In C++11 we renamed rvalue to prvalue.

Expressions




- In C++11 we renamed rvalue to prvalue.
- And we added a new value category: xvalue.

Expressions




Expressions




- A **glvalue** has a distinct address in memory.
 - i.e. it has an identity.

Expressions



Expressions



- Only rvalues will bind to an rvalue reference.
- lvalues will not bind to an rvalue reference.

Binding

```
void f(A& i, A j, A&& k);
```

Binding

Binds to lvalues



```
void f(A& i, A j, A&& k);
```

Binding

Binds to lvalues

```
void f(A& i, A j, A&& k);
```

Special case: Will
bind to rvalue if
const A&

Binding

Binds to lvalues
and rvalues

Binds to lvalues

```
void f(A& i, A j, A&& k);
```

Special case: Will
bind to rvalue if
const A&

Binding

Binds to lvalues
and rvalues

Binds to lvalues

```
void f(A& i, A j, A&& k);
```

Special case: Will
bind to rvalue if
const A&

lvalues require copy
rvalues require move
(the move can be elided)

Binding

Binds to lvalues
and rvalues

Binds to lvalues

Binds to rvalues

```
void f(A& i, A j, A&& k);
```

Special case: Will
bind to rvalue if
const A&

lvalues require copy
rvalues require move
(the move can be elided)

Types & Expressions

```
void f(A& i, A j, A&& k)
{
}
}
```

Types & Expressions

```
void f(A& i, A j, A&& k)
{
    ...
}
```

- i is declared as type A&.

Types & Expressions

```
void f(A& i, A j, A&& k)
{
    i; // lvalue A
}
```

- i is declared as type A&.
- The **expression** i has type A and is an lvalue.

Types & Expressions

```
void f(A& i, A j, A&& k)
{
    i; // lvalue A
}
```

- j is declared as type A.

Types & Expressions

```
void f(A& i, A j, A&& k)
{
    i; // lvalue A
    j; // lvalue A
}
```

- j is declared as type A.
- The **expression** j has type A and is an lvalue.

Types & Expressions

```
void f(A& i, A j, A&& k)
{
    i; // lvalue A
    j; // lvalue A
}
```

- k is declared as type A&&.

Types & Expressions

```
void f(A& i, A j, A&& k)
{
    i; // lvalue A
    j; // lvalue A
    k; // lvalue A
}
```

- k is declared as type A&&.
- The **expression** k has type A and is an lvalue.

Types & Expressions

```
void f(A& i, A j, A&& k)
{
}
}
```

Types & Expressions

```
void g(A&);    // #1
void g(A&&);  // #2
void f(A& i, A j, A&& k)
{
}
}
```

Types & Expressions

```
void g(A&);    // #1
void g(A&&);  // #2
void f(A& i, A j, A&& k)
{
    g(i); // calls #1
}
```

Types & Expressions

```
void g(A&);    // #1
void g(A&&);  // #2
void f(A& i, A j, A&& k)
{
    g(i); // calls #1
    g(j); // calls #1
}
}
```

Types & Expressions

```
void g(A&);    // #1
void g(A&&);   // #2
void f(A& i, A j, A&& k)
{
    g(i); // calls #1
    g(j); // calls #1
    g(k); // calls #1
}
```

- The expression `k` is an lvalue `A`

Types & Expressions

```
void g(A&);    // #1
void g(A&&);   // #2
void f(A& i, A j, A&& k)
{
    g(static_cast<A&&>(i)); // calls #2
    g(static_cast<A&&>(j)); // calls #2
    g(static_cast<A&&>(k)); // calls #2
}
```

- An lvalue expression can be cast to an rvalue (xvalue) expression

Types & Expressions

```
void g(A&);    // #1
void g(A&&);   // #2
void f(A& i, A j, A&& k)
{
    g(std::move(i)); // calls #2
    g(std::move(j)); // calls #2
    g(std::move(k)); // calls #2
}
```

- Use `std::move` to perform the cast for better readability.

Outline

- The rvalue reference
- Move Semantics
- Factory Functions
- More rvalue ref rules
- “Perfect” forwarding

Outline

- The rvalue reference
- Move Semantics
- Factory Functions
- More rvalue ref rules
- “Perfect” forwarding

Observation

```
void f(A& i, A j, A&& k)
{
    // i is not a unique reference
    // j is a unique reference
    // k is a reference to xvalue or prvalue
}
```

Observation

```
void f(A& i, A j, A&& k)
{
    // i is not a unique reference
    // j is a unique reference
    // k is a reference to xvalue or prvalue
}
```

- `f()` can do anything it wants to `j`, as long as the object remains destructible.

Observation

```
void f(A& i, A j, A&& k)
{
    // i is not a unique reference
    // j is a unique reference
    // k is a reference to xvalue or prvalue
}
```

- `f()` can do anything it wants to `j`, as long as the object remains destructible.
- `f()` can do anything it wants to `k`, as long as `k` references a prvalue.

Observation

```
void f(A& i, A j, A&& k)
{
    // i is not a unique reference
    // j is a unique reference
    // k is a reference to xvalue or prvalue
}
```

- `f()` can do anything it wants to `j`, as long as the object remains destructible.
- `f()` can do anything it wants to `k`, as long as `k` references a prvalue.
- Convention: Do not cast an lvalue to an xvalue unless you want that object to be treated as a prvalue.

The move constructor

```
class A
{
    int* data_;          // heap allocated
public:
    A(const A& a);    // copy constructor
};


```

The move constructor

```
class A
{
    int* data_;          // heap allocated
public:
    A(const A& a);    // copy constructor
};


```

- copy constructor binds to an lvalue and copies resources.

The move constructor

```
class A
{
    int* data_;          // heap allocated
public:
    A(const A& a);    // copy constructor
    A(A&& a) noexcept // move constructor
        : data_(a.data_)
    { a.data_ = nullptr; }

};
```

- copy constructor binds to an lvalue and copies resources.
- move constructor binds to a rvalue and pilfers resources.

The move constructor

```
A make_A();
```

```
A a1;
```

```
A a2 = a1; // Calls copy ctor
```

The move constructor

```
A make_A();
```

```
A a1;
```

```
A a2 = a1; // Calls copy ctor
```

```
A a3 = make_A(); // Calls (or elides)  
// move ctor
```

The move constructor

```
A make_A();  
  
A a1;  
A a2 = a1;           // Calls copy ctor  
A a3 = make_A();    // Calls (or elides)  
                    //          move ctor  
A a4 = std::move(a1); // Calls move ctor
```

The move constructor

```
A make_A();  
  
A a1;  
A a2 = a1;           // Calls copy ctor  
A a3 = make_A();    // Calls (or elides)  
                    //          move ctor  
A a4 = std::move(a1); // Calls move ctor
```

- “Copies” from rvalues are made with the move constructor, which does nothing but trade pointers. *Fast!*

The move constructor

The move constructor

- If a class does not have a move constructor, its copy constructor will be used to copy from rvalues (just as in C++98/03).

The move constructor

- If a class does not have a move constructor, its copy constructor will be used to copy from rvalues (just as in C++98/03).
- Scalars move the same as they copy.

The move constructor

The move constructor

```
struct A
{
    A(const A& a);
    A(A&&) = default;
};
```

- Copy and move constructors can be explicitly defaulted.
- The default copies/moves each base and data member (unless it is defined as deleted).

The move constructor

The move constructor

```
struct member
{
    member(const member&);

};

struct A
{
    member m_;
    A(A&&) = default; // deleted
};
```

- A defaulted move constructor is defined as deleted if:
 - there is a base or member with no move constructor and it is not trivially copyable.

The move constructor

```
struct member
{
    member(const member&) = default;

};

struct A
{
    member m_;
    A(A&&) = default;
};
```

- A defaulted move constructor is defined as deleted if:
 - there is a base or member with no move constructor and it is not trivially copyable.

The move constructor

```
struct member
{
    member(const member&);
    member(member&&);

};

struct A
{
    member m_;
    A(A&&) = default;
};
```

- A defaulted move constructor is defined as deleted if:
 - there is a base or member with no move constructor and it is not trivially copyable.

The move constructor

```
struct member
{
    member(const member&);

};

struct A
{
    member m_;
    A(A&&) = default;
};
```

The move constructor

```
struct member
{
    member(const member&);

};

struct A
{
    member m_;
    A(A&&) = default;
};
```

- CWG issue 1402 (ready) changes the rules such that the defaulted move members will **not** be implicitly deleted, but instead copy the bases and members.

The move constructor

The move constructor

```
static_assert
(
    std::is_move_constructible<A>::value,
    "A should be move constructible"
);
```

- You can always test at compile time if a complete type is move constructible.

The move constructor

```
static_assert
(
    std::is_move_constructible<A>::value,
    "A should be move constructible"
);
```

- You can always test at compile time if a complete type is move constructible.
- This tests whether or not A is constructible from an rvalue A, not if A has a move constructor.
- But a type with a deleted move constructor is never move constructible.

The move constructor

The move constructor

```
struct A
{
    A(const A& a) = delete;
    A(A&&)      = default;
};
```

- Copy and move constructors can be explicitly deleted.

The move constructor

```
struct A
{
    A(const A& a) = default;
    A(A&&)      = delete;
};
```

- Copy and move constructors can be explicitly deleted.
- A deleted move constructor will prohibit copying from rvalues (rarely a good idea). Normally omit rather than delete a move constructor.

The move constructor

```
struct A
{
    A(const A& a) = default;

};
```

- Copy and move constructors can be explicitly deleted.
- A deleted move constructor will prohibit copying from rvalues (rarely a good idea). Normally omit rather than delete a move constructor.

The move constructor

The move constructor

```
struct A
{
    // A(const A&) = delete;
    // A& operator=(const A&) = delete;
    A(A&&);
};
```

- A user-declared move constructor (defaulted or not) will implicitly create a deleted copy constructor and copy assignment.

Implicit Special Members

```
class A
{
    std::string s_;
public:
    // A() noexcept = default;
    // A(const A&) noexcept = default;
    // A& operator=(const A&) noexcept = default;
    // A(A&&) noexcept = default;
    // A& operator=(A&&) noexcept = default;
    // ~A() noexcept = default;
};
```

noexcept is extension

- Comments indicate compiler supplied definitions.

Implicit Special Members

```
class A
{
    std::string s_;
public:
    A();
    // A(const A&) = default;
    // A& operator=(const A&) = default;
    // A(A&&) noexcept = default;
    // A& operator=(A&&) noexcept = default;
    // ~A() noexcept = default;
};
```

- Comments indicate compiler supplied definitions.

Implicit Special Members

```
class A
{
    std::string s_;
public:
    A(const A&);                                // deprecated
    // A& operator=(const A&) = default;
    // ~A() noexcept = default;
};
```

- Comments indicate compiler supplied definitions.

Implicit Special Members

```
class A
{
    std::string s_;
public:
    // A() noexcept = default;
    // A(const A&) = default;
    A& operator=(const A&) = default;

    // ~A() noexcept = default;
};
```

noexcept is extension
deprecated

- Comments indicate compiler supplied definitions.

Implicit Special Members

```
class A
{
    std::string s_;
public:
    // A() noexcept = default;
    // A(const A&) = default;
    // A& operator=(const A&) = default;
    ~A();
};
```

noexcept is extension
deprecated

- Comments indicate compiler supplied definitions.

Implicit Special Members

```
class A
{
    std::string s_;
public:

    // A(const A&) = delete;
    // A& operator=(const A&) = delete;
    A(A&&);

    // ~A() noexcept = default;
};
```

- Comments indicate compiler supplied definitions.

Implicit Special Members

```
class A
{
    std::string s_;
public:
    // A() noexcept = default;
    // A(const A&) = delete;
    // A& operator=(const A&) = delete;

    A& operator=(A&&);
    // ~A() noexcept = default;
};
```

noexcept is extension

- Comments indicate compiler supplied definitions.

Advice

- Put these (or other appropriate) tests right into your release code:

```
struct A
{
    std::string s_;
    std::vector<int> v_;

};
```

Advice

- Put these (or other appropriate) tests right into your release code:

```
struct A
{
    std::string s_;
    std::vector<int> v_;

};

// Howard says put these tests in!
static_assert(std::is_nothrow_default_constructible<A>::value, "");
static_assert(std::is_copy_constructible<A>::value, "");
static_assert(std::is_copy_assignable<A>::value, "");
static_assert(std::is_nothrow_move_constructible<A>::value, "");
static_assert(std::is_nothrow_move_assignable<A>::value, "");
static_assert(std::is_nothrow_destructible<A>::value, "");
```

Advice

- Put these (or other appropriate) tests right into your release code:

```
struct A
{
    std::string s_;
    std::vector<int> v_;
    A(const A&) = default;
};
```

// Howard says put these tests in!

```
static_assert(std::is_nothrow_default_constructible<A>::value, "");
static_assert(std::is_copy_constructible<A>::value, "");
static_assert(std::is_copy_assignable<A>::value, "");
static_assert(std::is_nothrow_move_constructible<A>::value, "");
static_assert(std::is_nothrow_move_assignable<A>::value, "");
static_assert(std::is_nothrow_destructible<A>::value, "");
```

Advice

- Put these (or other appropriate) tests right into your release code:

```
struct A
{
    std::string s_;
    std::vector<int> v_;
    A(const A&) = default;
};
```

```
// Howard says put these tests in! Or else!!!
```

```
static_assert(std::is_copy_constructible<A>::value, "");
static_assert(std::is_copy_assignable<A>::value, "");
```

```
static_assert(std::is_nothrow_destructible<A>::value, "");
```

The move assignment operator

The move assignment operator

- Everything that's been said about the move constructor applies to the move assignment operator.

The move assignment operator

```
class A
{
    int* data_;           // heap allocated
public:
    A& operator=(const A& a); // copy
};


```

The move assignment operator

```
class A
{
    int* data_;          // heap allocated
public:
    A& operator=(const A& a); // copy
```

- }; • copy assignment binds to lvalue rhs and copies resources.

The move assignment operator

```
class A
{
    int* data_;           // heap allocated
public:
    A& operator=(const A& a);    // copy
    A& operator=(A&& a) noexcept // move
    {
        std::swap(data_, a.data_);
        return *this;
    }
};
```

- copy assignment binds to lvalue rhs and copies resources.
- move assignment binds to rvalue rhs and does whatever is fastest to assume value of rhs.

The move assignment operator

The move assignment operator

```
class A
{
    fstream f_;
public:
    A& operator=(A&& a) noexcept
    {
        f_ = std::move(a.f_);
        return *this;
    }
};
```

- If your type holds std::lib components, move assigning those data members will generally do the right thing.

The move assignment operator

```
class A
{
    fstream f_;
public:
    A& operator=(A&& a) = default;
};


```

- If all you need to do is move assign bases and members, consider doing it with “= default”.

The move assignment operator

```
class A
{
    fstream f_;
public:
```

```
};
```

- Or doing it implicitly.

The move assignment operator

```
template <class T>
class A {
    T* data_;           // heap allocated
public:
    A& operator=(A&& a) noexcept {
        delete data_;
        data_ = a.data_;
        a.data_ = nullptr;
        return *this;
    }
};
```

- Does the move assignment operator need to check for self-assignment?

The move assignment operator

```
A& operator=(A&& a) noexcept {  
    delete data_;  
    data_ = a.data_;  
    a.data_ = nullptr;  
    return *this;  
}
```

The move assignment operator

```
A& operator=(A&& a) noexcept {  
    delete data_;  
    data_ = a.data_;  
    a.data_ = nullptr;  
    return *this;  
}
```

- Convention: Do not cast an lvalue to an xvalue unless you want that object to be treated as a prvalue.

The move assignment operator

```
A& operator=(A&& a) noexcept {  
    delete data_;  
    data_ = a.data_;  
    a.data_ = nullptr;  
    return *this;  
}
```

- Convention: Do not cast an lvalue to an xvalue unless you want that object to be treated as a prvalue.
- If ‘a’ refers to a prvalue, then it is not possible for ‘this’ and ‘a’ to refer to the same object.

The move assignment operator

```
A& operator=(A&& a) noexcept {  
    delete data_;  
    data_ = a.data_;  
    a.data_ = nullptr;  
    return *this;  
}
```

The move assignment operator

```
A& operator=(A&& a) noexcept {  
    delete data_;  
    data_ = a.data_;  
    a.data_ = nullptr;  
    return *this;  
}  
  
a = std::move(a);
```

- However if ‘a’ refers to an **xvalue**, then it is possible for ‘this’ and ‘a’ to refer to the same object.

The move assignment operator

```
A& operator=(A&& a) noexcept {  
    delete data_;  
    data_ = a.data_;  
    a.data_ = nullptr;  
    return *this;  
}  
  
a = std::move(a);
```

- However if ‘a’ refers to an **xvalue**, then it is possible for ‘this’ and ‘a’ to refer to the same object.
- But you’ve arguably broken convention.

The move assignment operator

```
A& operator=(A&& a) noexcept {
    delete data_;
    data_ = a.data_;
    a.data_ = nullptr;
    return *this;
}
```

The move assignment operator

```
A& operator=(A&& a) noexcept {  
    delete data_;  
    data_ = a.data_;  
    a.data_ = nullptr;  
    return *this;  
}
```

- Self-swap is one place where this can happen.

```
template <class T>  
void swap(T& x, T& y) {  
    T tmp(std::move(x));  
    x = std::move(y);  
    y = std::move(tmp);  
}  
std::swap(a, a);
```

The move assignment operator

```
A& operator=(A&& a) noexcept {
    delete data_;
    data_ = a.data_;
    a.data_ = nullptr;
    return *this;
}
template <class T>
void swap(T& x, T& y) {
    T tmp(std::move(x));
    x = std::move(y);
    y = std::move(tmp);
}
std::swap(a, a);
```

The move assignment operator

```
A& operator=(A&& a) noexcept {  
    delete data_;  
    data_ = a.data_;  
    a.data_ = nullptr;  
    return *this;  
}
```

- However in this case, the self-move assignment happens only on a moved-from value.

```
template <class T>  
void swap(T& x, T& y) {  
    T tmp(std::move(x));  
    x = std::move(y);  
    y = std::move(tmp);  
}  
std::swap(a, a);
```

The move assignment operator

```
A& operator=(A&& a) noexcept {
    delete data_;
    data_ = a.data_;
    a.data_ = nullptr;
    return *this;
}
template <class T>
void swap(T& x, T& y) {
    T tmp(std::move(x));
    x = std::move(y);
    y = std::move(tmp);
}
std::swap(a, a);
```

The move assignment operator

```
A& operator=(A&& a) noexcept {  
    delete data_;  
    data_ = a.data_;  
    a.data_ = nullptr;  
    return *this;  
}
```

- Self-move assignment from a moved-from value is most often naturally safe.

```
template <class T>  
void swap(T& x, T& y) {  
    T tmp(std::move(x));  
    x = std::move(y);  
    y = std::move(tmp);  
}  
  
std::swap(a, a);
```

The move assignment operator

```
A& operator=(A&& a) noexcept {  
    delete data_;  
    data_ = a.data_;  
    a.data_ = nullptr;  
    return *this;  
}
```

The move assignment operator

```
A& operator=(A&& a) noexcept {  
    delete data_;  
    data_ = a.data_;  
    a.data_ = nullptr;  
    return *this;  
}
```

- Indeed, in all permutation rearrangement algorithms (those that do not “remove” elements), the target of a move assignment is always in a “moved-from” state.

The move assignment operator

The move assignment operator

- I treat self-move assignment and self-swap as a performance bug in my code.

The move assignment operator

- I treat self-move assignment and self-swap as a performance bug in my code.
- My move assignment operators are “self-safe” only when in a moved-from state.

The move assignment operator

- I treat self-move assignment and self-swap as a performance bug in my code.
- My move assignment operators are “self-safe” only when in a moved-from state.
- Feel free to check for self-move assignment in your code.

The move assignment operator

- I treat self-move assignment and self-swap as a performance bug in my code.
- My move assignment operators are “self-safe” only when in a moved-from state.
- Feel free to check for self-move assignment in your code.
- Either with an if - to ignore the bug.

The move assignment operator

- I treat self-move assignment and self-swap as a performance bug in my code.
- My move assignment operators are “self-safe” only when in a moved-from state.
- Feel free to check for self-move assignment in your code.
 - Either with an if - to ignore the bug.
 - Or with an assert - to catch the bug.

Assignment boo boo's

```
class A
{
    std::vector<int> v_;
    std::string s_;
public:
    A& operator=(A a) {
        swap(a);
        return *this;
    }
};
```

- In C++98/03 it became popular to define assignment using a copy/swap pattern.
- This is very good if you need strong exception safety.

Assignment boo boo's

```
class A
{
    std::vector<int> v_;
    std::string s_;
public:
    A& operator=(A a) {
        swap(a);
        return *this;
    }
};
```

Assignment boo boo's

```
class A
{
    std::vector<int> v_;
    std::string s_;
public:
    A& operator=(A a) {
        swap(a);
        return *this;
    }
};
```

- It is very efficient when assigning from rvalues.

Assignment boo boo's

```
class A
{
    std::vector<int> v_;
    std::string s_;
public:
    A& operator=(A a) {
        swap(a);
        return *this;
    }
};
```

- It is very efficient when assigning from rvalues.
- It is not so efficient when assigning from lvalues.

Assignment boo boo's

```
class A
{
    std::vector<int> v_;
    std::string s_;
public:
    A& operator=(A a) {
        swap(a);
        return *this;
    }
};
```

Assignment boo boo's

```
class A
{
    std::vector<int> v_;
    std::string s_;
public:
    A& operator=(A a) {
        swap(a);
        return *this;
    }
};
```

- Strong exception safety is good, but it is not free.