

C++11 Smart Pointers and Algorithms

Howard Hinnant
May 16, 2012

Outline

- unique_ptr
- shared_ptr
- algorithms

Outline

- unique_ptr
- shared_ptr
- algorithms

unique_ptr

unique_ptr

- It is just like auto_ptr.

unique_ptr

- It is just like auto_ptr.
- But better.

unique_ptr

- Almost anything you can do with `auto_ptr`, you can do with `unique_ptr` using the same syntax:

```
auto_ptr<int>
factory(int i)
{
    return auto_ptr<int>(new int(i));
}
```

unique_ptr

- Almost anything you can do with `auto_ptr`, you can do with `unique_ptr` using the same syntax:

```
unique_ptr<int>
factory(int i)
{
    return unique_ptr<int>(new int(i));
}
```

unique_ptr

- Almost anything you can do with `auto_ptr`, you can do with `unique_ptr` using the same syntax:

```
void client(auto_ptr<int> p)
{
    // Ownership transferred into client
} // int* deleted here
```

unique_ptr

- Almost anything you can do with `auto_ptr`, you can do with `unique_ptr` using the same syntax:

```
void client(unique_ptr<int> p)
{
    // Ownership transferred into client
} // int* deleted here
```

unique_ptr

- Almost anything you can do with `auto_ptr`, you can do with `unique_ptr` using the same syntax:

```
void test()
{
    auto_ptr<int> p = factory(2);
    // *p == 2
    p.reset(new int(3));
    // *p == 3
    client(factory(4));
}
```

unique_ptr

- Almost anything you can do with `auto_ptr`, you can do with `unique_ptr` using the same syntax:

```
void test()
{
    unique_ptr<int> p = factory(2);
    // *p == 2
    p.reset(new int(3));
    // *p == 3
    client(factory(4));
}
```

unique_ptr

- Ok, so what is the point of unique_ptr?

unique_ptr

- It is both easy and common to write generic code that looks like this:

unique_ptr

- It is both easy and common to write generic code that looks like this:

```
template <class T>
void foo(T t)
{
    T copy_of_t = t;
    assert(copy_of_t == t);
}
```



This assert is most often implicit in the code logic (not a literal explicit assert).

unique_ptr

- It is both easy and common to write generic code that looks like this:
- Early implementations of sort did just this by copying the pivot element from the sequence, with the subsequent logic assuming this was a copy.

```
template <class I>
void sort(I first, I last)
{
    // ...
    value_type pivot = *middle;
    // ...
}
```

unique_ptr

- It is both easy and common to write generic code that looks like this:
- Early implementations of sort did just this by copying the pivot element from the sequence, with the subsequent logic assuming this was a copy.
 - Sorting sequences of auto_ptr subsequently failed at run time because the expression that looked like a copy was really a move.

unique_ptr

unique_ptr

- In order to be safely usable in generic code, copying must have copy syntax, and moving must have some **other** syntax.

unique_ptr

- In order to be safely usable in generic code, copying must have **copy syntax**, and moving must have some **other syntax**.
- `auto_ptr` is unsafe because it moves with **copy syntax**. It is now deprecated.

unique_ptr

- In order to be safely usable in generic code, copying must have **copy syntax**, and moving must have some **other syntax**.
- `auto_ptr` is unsafe because it moves with **copy syntax**. It is now deprecated.
- `unique_ptr` will not compile if **copy syntax** is used. But it can be moved with **syntax** that can not be mistaken for a **copy**.

unique_ptr

- `unique_ptr` is a “move-only” type.
- It can not be copied, but it can be moved.

unique_ptr

- unique_ptr is a “move-only” type.
- It can not be copied, but it can be moved.

```
unique_ptr<int> p1(new int(3));
```

```
unique_ptr<int> p2 = p1; // Does not compile!
```

unique_ptr

- unique_ptr is a “move-only” type.
- It can not be copied, but it can be moved.

```
unique_ptr<int> p1(new int(3));
```

```
unique_ptr<int> p2 = std::move(p1); // Ok!
```

unique_ptr

- unique_ptr is a “move-only” type.
- It can not be copied, but it can be moved.

```
unique_ptr<int> p1(new int(3));
```

```
unique_ptr<int> p2 = source(); // Also Ok!
```

unique_ptr

- `unique_ptr<Derived>` can convert to `unique_ptr<Base>`.

unique_ptr

- `unique_ptr<Derived>` can convert to `unique_ptr<Base>`.

```
unique_ptr<Derived> source(); // function
```

```
unique_ptr<Base> p = source();
```

`~Base()` must be virtual

unique_ptr

- `unique_ptr` has a *custom deleter*.

unique_ptr

- unique_ptr has a custom deleter.

```
struct close_stream
{
    void operator()(std::ofstream* os) const
    {os->close();}
};

typedef unique_ptr<std::ofstream, close_stream>
FilePtr;
```

point to this → and call this at destruction

which does this ↑

unique_ptr

- `unique_ptr` has a *custom deleter*.

```
typedef unique_ptr<std::ofstream, close_stream>
    FilePtr;
```

unique_ptr

- `unique_ptr` has a *custom deleter*.

```
typedef unique_ptr<std::ofstream, close_stream>  
    FilePtr;
```

- `FilePtr` owns the open state of an `ofstream`, not the object itself.

unique_ptr

- `unique_ptr` has a *custom deleter*.

```
typedef unique_ptr<std::ofstream, close_stream>  
FilePtr;
```

unique_ptr

- `unique_ptr` has a *custom deleter*.

```
typedef unique_ptr<std::ofstream, close_stream>
    FilePtr;

FilePtr get_log()
{
    static std::ofstream log_file;
    log_file.open("log file");
    return FilePtr(&log_file);
}

void foo()
{
    FilePtr fp = get_log();
    *fp << "some text\n";
} // fp->close()
```

unique_ptr

- unique_ptr has a *custom deleter*.

unique_ptr

- `unique_ptr` has a *custom deleter*.
- If the deleter is “empty”, space for it will be optimized away.
 - The default deleters are empty.
 - `sizeof(unique_ptr<T>) == sizeof(T*)`

unique_ptr

- `unique_ptr` has a *custom deleter*.

unique_ptr

- `unique_ptr` has a *custom deleter*.
- A deleter can be an `lvalue reference type`.
- This can be useful if you want to keep the state for a deleter in one place.

```
unique_ptr<T, MyDeleter&> p(ptr, deleter);
```

unique_ptr

- `unique_ptr` has a *custom deleter*.
- A deleter can be a function pointer.

unique_ptr

- unique_ptr has a custom deleter.
- A deleter can be a function pointer.

```
template <class T>
unique_ptr<char, void(*)(void*)> type_name(const T&)
{
    return unique_ptr<char, void(*)(void*)>(
        __cxa_demangle(typeid(T).name(), nullptr,
                       nullptr, nullptr),
        free);
}
```

This returns
malloc'd memory
memory.

This will be used to
deallocate the memory.

unique_ptr

- `unique_ptr` has a *custom deleter*.

```
template <class T>
unique_ptr<char, void(*)(void*)>
type_name(const T&)
{
    return unique_ptr<char, void(*)(void*)>
    (
        __cxa_demangle(typeid(T).name(), nullptr,
                       nullptr, nullptr),
        free
    );
}
```

unique_ptr

- `unique_ptr` has a *custom deleter*.

```
template <class T>
unique_ptr<char, void(*)(void*)>
type_name(const T&)
{
    return unique_ptr<char, void(*)(void*)>
    (
        __cxa_demangle(typeid(T).name(), nullptr,
                       nullptr, nullptr),
        free
    );
}

cout << type_name(x).get() << '\n';
```

unique_ptr

```
cout << type_name(x).get() << '\n';
```

- compare to:

```
char* name = nullptr;
try
{
    name = __cxa_demangle(typeid(x).name(), nullptr,
                           nullptr, nullptr);
    cout << name << '\n';
    free(name);
}
catch (...)
{
    free(name);
    throw;
}
```

unique_ptr

- `unique_ptr` has an array form:

unique_ptr

- `unique_ptr` has an array form:

`unique_ptr<T[], D>`

unique_ptr

- unique_ptr has an array form:

`unique_ptr<T[], D>`

- This unique_ptr does not have operator*().

unique_ptr

- unique_ptr has an array form:

`unique_ptr<T[], D>`

- This unique_ptr does not have operator*().
- Nor will it convert from derived to base.

unique_ptr

- unique_ptr has an array form:

`unique_ptr<T[], D>`

- This unique_ptr does not have operator*().
- Nor will it convert from derived to base.
- But it does have `T& operator[](size_t) const.`

unique_ptr

- unique_ptr has an array form:

unique_ptr<T[], D>

- This unique_ptr does not have operator*().
- Nor will it convert from derived to base.
- But it does have T& operator[](size_t) const.
- The default deleter uses delete[].

```
unique_ptr<char[]> p(new char[10]);  
p[0] = 'a';
```

unique_ptr

- unique_ptr has an array form:

```
template <class T>
unique_ptr<char[], void(*)(void*)>
type_name(const T&)
{
    return unique_ptr<char[], void(*)(void*)>
    (
        __cxa_demangle(typeid(T).name(), nullptr,
                       nullptr, nullptr),
        free
    );
}
```

unique_ptr

braces
added

- unique_ptr has an array form:

```
template <class T>
unique_ptr<char[], void(*)(void*)>
type_name(const T&)
{
    return unique_ptr<char[], void(*)(void*)>
    (
        __cxa_demangle(typeid(T).name(), nullptr,
                       nullptr, nullptr),
        free
    );
}
auto nm = type_name(x);
for (int i = 0; nm[i]; ++i)
    cout << nm[i];
```

indexing
can now
be used

deallocation still uses

free

unique_ptr

- `unique_ptr` has support for incomplete types (useful for pimpl).

unique_ptr

- unique_ptr has support for incomplete types (useful for pimpl).

```
class A
{
    class impl;
    unique_ptr<impl> ptr_;
public:
    A();
    A(A&&);
    A& operator=(A&&);
    ~A();
};
```

unique_ptr

- unique_ptr has support for incomplete types (useful for pimpl).

```
class A::impl {};  
  
A::A() = default;  
A::A(A&&) = default;  
A& A::operator=(A&&) = default;  
A::~A() = default;
```

unique_ptr

- unique_ptr has support for incomplete types (useful for pimpl).

```
class A::impl {};
```

```
A::A() = default;  
A::A(A&&) = default;  
A& A::operator=(A&&) = default;  
A::~A() = default;
```

- Each special member can be defaulted, but must be outlined into a source once A::impl is complete.

unique_ptr

- unique_ptr has support for incomplete types (useful for pimpl).
- If you accidentally attempt to do anything with the incomplete A::impl that is not allowed, a compile-time error is guaranteed.

unique_ptr

- `unique_ptr` has support for custom storage (internal pointer type).

unique_ptr

- unique_ptr has support for custom storage (internal pointer type).
- Use this to put unique_ptr into process-shared memory, using an “offset_ptr” as the storage.

unique_ptr

- `unique_ptr` has support for custom storage (internal pointer type).

unique_ptr

- `unique_ptr` has support for custom storage (internal pointer type).

```
template <class T>
struct MyDeleter
{
    struct pointer
    {
        // emulate a pointer ...
    };

    void operator()(pointer p);
};

unique_ptr<int, MyDeleter<int>> p;
```

unique_ptr

- `unique_ptr` has support for custom storage (internal pointer type).

```
unique_ptr<int, MyDeleter<int>> p;
```

unique_ptr

- `unique_ptr` has support for custom storage (internal pointer type).

```
unique_ptr<int, MyDeleter<int>> p;
```

`unique_ptr<int, MyDeleter<int>>::pointer`
is the same type as
`MyDeleter<int>::pointer`

unique_ptr

- `unique_ptr` has support for custom storage (internal pointer type).

```
unique_ptr<int, MyDeleter<int>> p;
```

`MyDeleter<int>::pointer`

```
unique_ptr<int, MyDeleter<int>::pointer
```

unique_ptr

- `unique_ptr` has support for custom storage (internal pointer type).

```
unique_ptr<int, MyDeleter<int>> p;
```

If

`MyDeleter<int>::pointer`

does not exist, then

```
unique_ptr<int, MyDeleter<int>::pointer
```

is

`int*`

unique_ptr

- `unique_ptr` can be put into containers and manipulated with algorithms.

unique_ptr

- `unique_ptr` can be put into containers and manipulated with algorithms.

```
int main() {  
    typedef unique_ptr<int> Ptr;  
    Ptr p[] = {Ptr(new int(2)), Ptr(new int(3)),  
              Ptr(new int(1))};
```

unique_ptr

- `unique_ptr` can be put into containers and manipulated with algorithms.

```
int main() {
    typedef unique_ptr<int> Ptr;
    Ptr p[] = {Ptr(new int(2)), Ptr(new int(3)),
               Ptr(new int(1))};
    vector<Ptr> v(make_move_iterator(begin(p)),
                   make_move_iterator(end(p)));
```

unique_ptr

- `unique_ptr` can be put into containers and manipulated with algorithms.

```
int main() {
    typedef unique_ptr<int> Ptr;
    Ptr p[] = {Ptr(new int(2)), Ptr(new int(3)),
               Ptr(new int(1))};
    vector<Ptr> v(make_move_iterator(begin(p)),
                   make_move_iterator(end(p)));
    sort(v.begin(), v.end(), [](const Ptr& x,
                                const Ptr& y)
         {return *x < *y;});
```

unique_ptr

- `unique_ptr` can be put into containers and manipulated with algorithms.

```
int main() {
    typedef unique_ptr<int> Ptr;
    Ptr p[] = {Ptr(new int(2)), Ptr(new int(3)),
               Ptr(new int(1))};
    vector<Ptr> v(make_move_iterator(begin(p)),
                   make_move_iterator(end(p)));
    sort(v.begin(), v.end(), [](const Ptr& x,
                                const Ptr& y)
         {return *x < *y;});
    for (const auto& p : v)
        cout << *p << ' ';
    cout << '\n';
}
```

unique_ptr

- Use `const unique_ptr<T>` when you want to guarantee that ownership is not transferred out of scope.

unique_ptr

- Use `const unique_ptr<T>` when you want to guarantee that ownership is not transferred out of scope.

```
const unique_ptr<Base> p(new Derived);  
// ...  
return p; // Compile-time error
```

unique_ptr


- Use `const unique_ptr<T>` when you want to guarantee that ownership is not transferred out of scope.

```
const unique_ptr<Base> p(new Derived);
// ...
swap(p, p2); // Compile-time error
```

unique_ptr

- Use `const unique_ptr<T>` when you want to guarantee that ownership is not transferred out of scope.
- `const std::unique_ptr<T>` is arguably a better type than `boost::scoped_ptr<T>` for guaranteeing that ownership is not transferred out of scope.

Migrating from auto_ptr to unique_ptr



Migrating from auto_ptr to unique_ptr

- It is safe to do a global search for “auto_ptr” and replace with “unique_ptr”.
- If the result compiles, you are good to go.

Migrating from auto_ptr to unique_ptr

- It is safe to do a global search for “auto_ptr” and replace with “unique_ptr”.
- If the result compiles, you are good to go.
- If the result does not compile, it probably looks something like this:

```
p1 = p2;
```

Migrating from auto_ptr to unique_ptr

- If the result does not compile, it probably looks something like this:

Migrating from auto_ptr to unique_ptr

- If the result does not compile, it probably looks something like this:
- Try this instead:

```
p1 = std::move(p2);
```

- But inspect the code to see if p2 is inappropriately used afterwards:

```
f(p2);
```

- Why is f() being called with a moved-from smart pointer?

Migrating from auto_ptr to unique_ptr

- People have confirmed to me that this exercise has found bugs in large projects.

Outline

- unique_ptr
- shared_ptr
- algorithms

Outline


- unique_ptr
- shared_ptr
- algorithms

shared_ptr

- The most bullet-proof, most functional reference counted pointer.
- This is the same shared_ptr you know and love from boost.

`shared_ptr`


A contrast with `unique_ptr`



shared_ptr

A contrast with unique_ptr


- shared_ptr models shared ownership.



shared_ptr


A contrast with unique_ptr

- shared_ptr models shared ownership.
- unique_ptr models unique ownership.



shared_ptr


A contrast with unique_ptr



shared_ptr


A contrast with unique_ptr

- A unique_ptr can be converted to a shared_ptr.
- But not vice-versa.



`shared_ptr`

A contrast with `unique_ptr`




`shared_ptr`

A contrast with `unique_ptr`

- `sizeof(shared_ptr<T>)` == 2 words
- `sizeof(unique_ptr<T>)` == 1 word

`shared_ptr`

A contrast with `unique_ptr`



`shared_ptr`

A contrast with `unique_ptr`

- Both `shared_ptr` and `unique_ptr` support custom deleters, but...
- `shared_ptr`'s deleter is not part of its type. It is only specified in a constructor.
- `unique_ptr`'s deleter is part of the type. One can optionally specify a deleter in the constructor.

`shared_ptr`

A contrast with `unique_ptr`

- Both `shared_ptr` and `unique_ptr` support custom deleters, but...

`shared_ptr`

A contrast with `unique_ptr`

- Both `shared_ptr` and `unique_ptr` support custom deleters, but...
- Clients of a `shared_ptr` factory function do not need to know about the type of the deleter the `shared_ptr` was constructed with.
- There is only one type of `shared_ptr`.

`shared_ptr`

A contrast with `unique_ptr`

- Both `shared_ptr` and `unique_ptr` support custom deleters, but...
- Clients of a `shared_ptr` factory function do not need to know about the type of the deleter the `shared_ptr` was constructed with.
 - There is only one type of `shared_ptr`.
 - Clients of a `unique_ptr` factory function must know the type of the `unique_ptr`'s deleter.

`shared_ptr`

A contrast with `unique_ptr`

- Both `shared_ptr` and `unique_ptr` support custom deleters, but...
- Some of the “hidden deleter” abstraction can be gained for `unique_ptr` by using a function pointer for the deleter, and hiding what function it points to:

shared_ptr

A contrast with unique_ptr

- Both shared_ptr and unique_ptr support custom deleters, but...
- Some of the “hidden deleter” abstraction can be gained for unique_ptr by using a function pointer for the deleter, and hiding what function it points to:

```
unique_ptr<T, void(*)(void*)>  
source();
```



Deallocated by std::free,
or maybe something else?

`shared_ptr`

A contrast with `unique_ptr`

- Both `shared_ptr` and `unique_ptr` support custom deleters, but...

`shared_ptr`

A contrast with `unique_ptr`

- Both `shared_ptr` and `unique_ptr` support custom deleters, but...
 - The `shared_ptr` “type-erased” deleter design requires that the deleter be stored on the heap.
 - `shared_ptr<T>(new T)` requires two allocations: one for the `T` and one for the control block holding the deleter.

`shared_ptr`

A contrast with `unique_ptr`

- Both `shared_ptr` and `unique_ptr` support custom deleters, but...
 - The `unique_ptr` deleter is stored in the pointer itself.
 - `unique_ptr<T>(new T)` requires only one allocation.

`shared_ptr`

A contrast with `unique_ptr`

- Both `shared_ptr` and `unique_ptr` support custom deleters, but...
 - Use of `make_shared<T>(Args...)` can reduce the number of allocations required to construct a `shared_ptr` down to one.
 - There is no `make_unique<T>(Args...)` at this time...

`shared_ptr`

A contrast with `unique_ptr`

- `shared_ptr<T[]>` is not supported at this time.
- But you can use a custom deleter to get the right deallocation:

shared_ptr

A contrast with unique_ptr

- `shared_ptr<T[]>` is not supported at this time.
- But you can use a custom deleter to get the right deallocation:

```
shared_ptr<T> p(new T[3],  
                 default_delete<T[]>());
```

`shared_ptr`


Ownership cycles

- Cyclic ownership is a constant danger when reference counted pointers are used indiscriminately.

shared_ptr

Ownership cycles

- Cyclic ownership is a constant danger when reference counted pointers are used indiscriminately.




shared_ptr

Ownership cycles

- Cyclic ownership is a constant danger when reference counted pointers are used indiscriminately.


A cycle can not be deleted because there is always a reference count keeping everything alive.



shared_ptr

Ownership cycles


- The solution is to use `weak_ptr` to break the cycle.
- `weak_ptr` does not own what it points to.
- `weak_ptr` knows when its pointee gets deleted.



shared_ptr

Ownership cycles

- The solution is to use `weak_ptr` to break the cycle.
- `weak_ptr` does not own what it points to.
- `weak_ptr` knows when its pointee gets deleted.



`shared_ptr`

Ownership cycles

- The solution is to use `weak_ptr` to break the cycle.
- `weak_ptr` does not own what it points to.
- `weak_ptr` knows when its pointee gets deleted.

shared_ptr

Ownership cycles

```
class C
{
    shared_ptr<A> ptr_;
public:
    C() = default;
    ~C() {std::cout << "~C()\n";}
    void set(shared_ptr<A> p) {ptr_ = p;}
};
```

shared_ptr

Ownership cycles

To fix, simply replace
shared_ptr with weak_ptr.

```
class C
{
    weak_ptr<A> ptr_;
public:
    C() = default;
    ~C() {std::cout << "~C()\n";}
    void set(shared_ptr<A> p) {ptr_ = p;}
};
```

shared_ptr

Ownership cycles

To fix, simply replace
shared_ptr with weak_ptr.

```
class C
{
public:
    C() = default;
    ~C() {std::cout << "~C()\n";}
    void set(shared_ptr<A> p) {ptr_ = p;}
};
```

`shared_ptr`

`weak_ptr`

- `weak_ptr` must be converted to `shared_ptr` in order to be dereferenced.
- Explicit construction will throw `bad_weak_ptr` if `expired`.

```
class C
{
    weak_ptr<A> ptr_;
public:
    shared_ptr<A> get() const
        {return shared_ptr<A>(ptr_);}
};
```

`shared_ptr`

`weak_ptr`

- `weak_ptr` must be converted to `shared_ptr` in order to be dereferenced.
- Use `lock()` to instead return a null `shared_ptr` when expired.

```
class C
{
    weak_ptr<A> ptr_;
public:
    shared_ptr<A> get() const
};
```

`shared_ptr`

`weak_ptr`

- `weak_ptr` must be converted to `shared_ptr` in order to be dereferenced.
- Use `lock()` to instead return a null `shared_ptr` when expired.

```
class C
{
    weak_ptr<A> ptr_;
public:
    shared_ptr<A> get() const
        {return ptr_.lock();}
};
```

`shared_ptr`

`weak_ptr`

- Not being able to directly dereference a `weak_ptr` is a critical safety feature in multithreaded code.

shared_ptr

weak_ptr

- Not being able to directly dereference a `weak_ptr` is a critical safety feature in multithreaded code.

Thread A

```
weak_ptr<A> wp = ...  
if ( !wp.expired() )  
    wp->do_something();
```

Thread B

```
sp.reset();
```



Won't even compile!

What `wp` refers to could get
destructed during `do_something()`!

shared_ptr

weak_ptr

- Not being able to directly dereference a `weak_ptr` is a critical safety feature in multithreaded code.

Thread A

```
weak_ptr<A> wp = ...  
shared_ptr<A> sp(wp);  
sp->do_something();
```

Thread B

```
sp.reset();
```

`shared_ptr`

`weak_ptr`

- Not being able to directly dereference a `weak_ptr` is a critical safety feature in multithreaded code.

Thread A

```
weak_ptr<A> wp = ...  
shared_ptr<A> sp(wp);  
sp->do_something();
```

Thread B

```
sp.reset();
```

A successful conversion to `shared_ptr`, atomically extends the life time long enough to complete `do_something()`!

Which Smart Pointer is Preferred?

- So which smart pointer should I reach for first?

Which Smart Pointer is Preferred?

- Neither!

Which Smart Pointer is Preferred?

- Prefer holding data members directly.

```
class A
{
    B b_;
};
```

- Use pointers, even smart ones, sparingly.

Which Smart Pointer is Preferred?

- Use a smart pointer when you need to point to a base class.

```
class A
{
    unique_ptr<Base> ptr_;
};
```

- Prefer unique_ptr to model a single owner.
- Unique ownership is simpler to reason about than shared ownership.

Which Smart Pointer is Preferred?

```
class A
{
    unique_ptr<Base> ptr_;
public:
    A(const A& a)
        : ptr_(a.ptr_ ?
                a.ptr_->clone() :
                nullptr) {}
    A(A&&) noexcept = default;
    // ...
};
```


- Do not choose `shared_ptr` just to enable cheap copies. Let move semantics take care of that.

Which Smart Pointer is Preferred?

- Do not choose `shared_ptr` just to enable cheap copies. Let move semantics take care of that.
 - Most copies turn into moves in C++11.
 - Moving a `unique_ptr` is twice as fast as moving a `shared_ptr`.

Which Smart Pointer is Preferred?

- Do not choose `shared_ptr` just to enable cheap copies. Let move semantics take care of that.
 - Most copies turn into moves in C++11.
 - Moving a `unique_ptr` is twice as fast as moving a `shared_ptr`.



Which Smart Pointer is Preferred?

```
class A
{
    shared_ptr<Base> ptr_;
public:
    // ...
};
```

- Choose `shared_ptr` only when you actually need shared ownership semantics, or when you know you will need copies (not moves) and the pointee is *always* immutable.

Outline

- unique_ptr
- shared_ptr
- algorithms

Outline

- unique_ptr
- shared_ptr
- algorithms

<algorithm>

- Many sequence-permuting standard algorithms will now use move or swap, not copy.



<algorithm>

- Many sequence-permuting standard algorithms will now use move or swap, not copy.



- These algorithms no longer even require copyability: they will work with move-only types such as `unique_ptr<T>`.

<algorithm>

- Many sequence-permuting standard algorithms will now use move or swap, not copy.

<algorithm>

- Many sequence-permuting standard algorithms will now use move or swap, not copy.

swap_ranges	random_shuffle	
iter_swap	sort	push_heap
remove	inplace_merge	
partition	remove_if	pop_heap
stable_partition	stable_sort	make_heap
next_permutation	prev_permutation	
reverse	partial_sort	
unique	rotate	sort_heap
	nth_element	

<algorithm>

- Many sequence-permuting standard algorithms will now use move or swap, not copy.

<algorithm>

- Many sequence-permuting standard algorithms will now use move or swap, not copy.

```
vector<unique_ptr<T>> v = ...  
sort(v.begin(), v.end(),  
     [](const unique_ptr<T>& x,  
        const unique_ptr<T>& y)  
     {return *x < *y;});
```

<algorithm>

- Many sequence-permuting standard algorithms will now use move or swap, not copy.

```
vector<unique_ptr<T>> v = ...  
rotate(v.begin(), v.begin() + 5, v.end());
```

<algorithm>

- Many sequence-permuting standard algorithms will now use move or swap, not copy.

```
vector<unique_ptr<T>> v = ...  
remove(v.begin(), v.end(), nullptr);
```

<algorithm>

- Many sequence-permuting standard algorithms will now use move or swap, not copy.

```
vector<unique_ptr<T>> v = ...  
remove(v.begin(), v.end(), nullptr);
```

Remove all of the nulls in
the sequence.

<algorithm>

- Many sequence-permuting standard algorithms will now use move or swap, not copy.

`vector<unique_ptr<T>> v = ...`

Consider using:

`std::vector<std::unique_ptr<T>>`
over:
`boost::ptr_vector<T>.`

<algorithm>

- Many sequence-permuting standard algorithms will now use move or swap, not copy.

```
vector<unique_ptr<T>> v = ...
```

Consider using:

std::vector<std::unique_ptr<T>>
over:
boost::ptr_vector<T>.

You get a wider range of available algorithms.

<algorithm>

New algorithms

- There are about 20 new algorithms in C++11

<algorithm>

New algorithms

- There are about 20 new algorithms in C++11
- Here are a few of my favorites...

<algorithm>

New algorithms

- `minmax_element`

<algorithm>

New algorithms

- `minmax_element`
 - Find both the minimum and maximum.

```
template<class ForwardIterator>
pair<ForwardIterator, ForwardIterator>
minmax_element(ForwardIterator first,
                ForwardIterator last);
```

```
template<class ForwardIterator,
         class Compare>
pair<ForwardIterator, ForwardIterator>
minmax_element(ForwardIterator first,
                ForwardIterator last,
                Compare comp);
```

<algorithm>

New algorithms

- `minmax_element`

<algorithm>

New algorithms

- `minmax_element`

```
struct A
{
    static int comp_count;
    int data;
    A(int d) : data(d) {}

    friend std::ostream&
    operator<<(std::ostream& os, const A& x)
    {return os << x.data;}
};
```

<algorithm>

New algorithms

- `minmax_element`

<algorithm>

New algorithms

- `minmax_element`

```
int A::comp_count = 0;  
  
bool  
operator <(const A& x, const A& y)  
{  
    ++A::comp_count;  
    return x.data < y.data;  
}
```

<algorithm>

New algorithms

- `minmax_element`

<algorithm>

New algorithms

- `minmax_element`

```
int main()
{
    mt19937_64 eng;
    uniform_int_distribution<> d(0, 1000000);
    vector<A> v;
    for (int i = 0; i < 1000; ++i)
        v.push_back(d(eng));
    auto p = minmax_element(v.begin(), v.end());
    cout << "v.size() = " << v.size() << '\n';
    cout << "min = " << *p.first << '\n';
    cout << "max = " << *p.second << '\n';
    cout << "# of comparisons = "
        << A::comp_count << '\n';
}
```

<algorithm>

New algorithms

- `minmax_element`

<algorithm>

New algorithms

- `minmax_element`

```
v.size() = 1000  
min = 629  
max = 999101  
# of comparisons = 1498
```

<algorithm>

New algorithms

- `minmax_element`

<algorithm>

New algorithms

- `minmax_element`

```
auto p =  
    std::make_pair  
(  
        std::min_element(v.begin(), v.end()),  
        std::max_element(v.begin(), v.end()))  
    );
```

<algorithm>

New algorithms

- `minmax_element`

<algorithm>

New algorithms

- `minmax_element`

```
minmax    v.size() = 1000  
          min = 629  
          max = 999101  
          # of comparisons = 1498
```

```
min + max  v.size() = 1000  
          min = 629  
          max = 999101  
          # of comparisons = 1998
```

`minmax` is at least 33% faster than using `min` and `max`!
And it is easier to use.

<algorithm>

New algorithms

- `is_sorted_until`

<algorithm>

New algorithms

- `is_sorted_until`
- Find sorted prefix of sequence.


```
template<class ForwardIterator>
ForwardIterator
is_sorted_until(ForwardIterator first,
               ForwardIterator last);
```

```
template <class ForwardIterator,
          class Compare>
ForwardIterator
is_sorted_until(ForwardIterator first,
               ForwardIterator last,
               Compare comp);
```

<algorithm>

New algorithms


- `is_sorted_until`
- Find sorted prefix of sequence.



<algorithm>

New algorithms

- `is_sorted_until`
- Find sorted prefix of sequence.



<algorithm>

New algorithms


- `is_sorted_until`
- Find sorted prefix of sequence.

<algorithm>

New algorithms

- `is_sorted_until`
 - Find sorted prefix of sequence.
 - Can be used for faster sorts when it is known that there is a large sorted prefix.

```
std::vector<int> v;
for (int i = 0; i < 100000; ++i)
    v.push_back(i);
for (int i = 0; i < 10000; ++i)
    v.push_back(d(eng));
```



uniform_int_distribution

<algorithm>

New algorithms

- `is_sorted_until`
- Find sorted prefix of sequence.

```
typedef std::chrono::high_resolution_clock Clock;
typedef std::chrono::duration<float, std::micro>
                           microsec;
auto t0 = Clock::now();
  
  
auto t1 = Clock::now();
std::cout << microsec(t1-t0).count() << " ms\n";
```

<algorithm>

New algorithms

- `is_sorted_until`
- Find sorted prefix of sequence.

```
typedef std::chrono::high_resolution_clock Clock;
typedef std::chrono::duration<float, std::micro>
                           microsec;
auto t0 = Clock::now();
auto t1 = Clock::now();
std::cout << microsec(t1-t0).count() << " ms\n";
```

High resolution timer boiler plate

The diagram consists of three main parts: two code snippets at the bottom and a descriptive text block above them. Two arrows point from the variable declarations 'Clock' and 'duration' in the first snippet to the word 'microsec' in the descriptive text. Another arrow points from the assignment 't0 = Clock::now();' to the first 'Clock::now()' in the second snippet.

<algorithm>

New algorithms

- `is_sorted_until`
- Find sorted prefix of sequence.

```
typedef std::chrono::high_resolution_clock Clock;
typedef std::chrono::duration<float, std::micro>
                           microsec;
auto t0 = Clock::now();
std::stable_sort(v.begin(), v.end());
auto t1 = Clock::now();
std::cout << microsec(t1-t0).count() << " ms\n";
```

<algorithm>

New algorithms

- `is_sorted_until`
- Find sorted prefix of sequence.

```
typedef std::chrono::high_resolution_clock Clock;
typedef std::chrono::duration<float, std::micro>
                           microsec;

auto t0 = Clock::now();
auto i = std::is_sorted_until(v.begin(), v.end());
std::stable_sort(i, v.end());
std::inplace_merge(v.begin(), i, v.end());
auto t1 = Clock::now();
std::cout << microsec(t1-t0).count() << " ms\n";
```

<algorithm>

New algorithms

- `is_sorted_until`
 - Find sorted prefix of sequence.

```
auto i = std::is_sorted_until(v.begin(), v.end());  
std::stable_sort(i, v.end());  
std::inplace_merge(v.begin(), i, v.end());
```

<algorithm>

New algorithms

- `is_sorted_until`
 - Find sorted prefix of sequence.

```
auto i = std::is_sorted_until(v.begin(), v.end());  
std::stable_sort(i, v.end());  
std::inplace_merge(v.begin(), i, v.end());
```

is 2.6* times faster than:

```
std::stable_sort(v.begin(), v.end());
```

<algorithm>

New algorithms

- `is_sorted_until`
- Find sorted prefix of sequence.

```
auto i = std::is_sorted_until(v.begin(), v.end());  
std::stable_sort(i, v.end());  
std::inplace_merge(v.begin(), i, v.end());
```

is 2.6* times faster than:

```
std::stable_sort(v.begin(), v.end());
```

(*your mileage may vary)

<algorithm>

New algorithms

- `is_permutation`
 - Determine if one sequence is a permutation of another.

```
template<class FwdItr1, class FwdItr2>
bool
```

```
    is_permutation(FwdItr1 first1, FwdItr1 last1,
                  FwdItr2 first2);
```

```
template<class FwdItr1, class FwdItr2,
         class BinaryPred>
```

```
bool
    is_permutation(FwdItr1 first1, FwdItr1 last1,
                  FwdItr2 first2, BinaryPred pred);
```

<algorithm>

New algorithms

- `is_permutation`
 - Determine if one sequence is a permutation of another.

```
template<class FwdItr1, class FwdItr2,  
         class BinaryPred>  
bool  
is_permutation(FwdItr1 first1, FwdItr1 last1,  
               FwdItr2 first2, BinaryPred pred);
```

- Reasonably efficient: Linear if `equal(first1, last1, first2)`.

<algorithm>

New algorithms

- `is_permutation`
 - Determine if one sequence is a permutation of another.

```
template<class FwdItr1, class FwdItr2,  
         class BinaryPred>  
bool  
is_permutation(FwdItr1 first1, FwdItr1 last1,  
               FwdItr2 first2, BinaryPred pred);
```

- Reasonably efficient: Finds “false” cases quickly.


<algorithm>

New algorithms

- `is_permutation`
 - Determine if one sequence is a permutation of another.

```
int x[] = {1, 2, 3, 4, 5};  
const unsigned N = sizeof(x) / sizeof(int);  
int y[N] = {2, 1, 3, 5, 4};  
int count = 0;  
bool b = std::is_permutation(x, x+N, y,  
    [&](int a, int b) -> bool  
    {++count; return a == b;});
```

Summary



Summary

- Use `unique_ptr` for unique ownership.
 - You can use it in containers and with algorithms.

Summary

- Use `unique_ptr` for unique ownership.
 - You can use it in containers and with algorithms.
- Use `shared_ptr` for shared ownership.
 - Don't use it just because you need to put it in a container or use it with algorithms.

Summary

- Use `unique_ptr` for unique ownership.
 - You can use it in containers and with algorithms.
 - Use `shared_ptr` for shared ownership.
 - Don't use it just because you need to put it in a container or use it with algorithms.
 - Learn about and take advantage of new algorithms to make your code faster.

